

# Computational Radiative Transfer

---

University of Bath

Candidate Number: 24367

March 26, 2025



UNIVERSITY OF  
**BATH**

---

# PH30110: Computational Astrophysics

Candidate Number: 24367

Department of Physics,  
University of Bath, Bath  
BA2 7AY, UK

Date Submitted: March 26, 2025

## Contents

<b>1</b>	<b>Question 1: Random Numbers from a Non-Flat Distribution</b>	<b>2</b>
<b>2</b>	<b>Question 2: Monte Carlo Scattering, Isotropic</b>	<b>3</b>
<b>3</b>	<b>Question 3: Rayleigh scattering</b>	<b>5</b>
<b>4</b>	<b>Appendix</b>	<b>9</b>
4.1	Full Code . . . . .	9

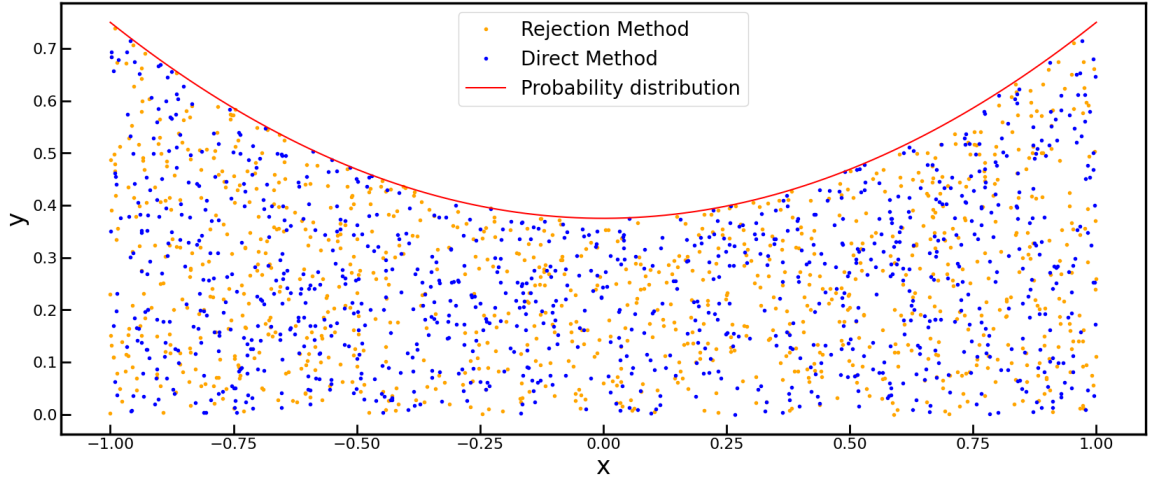


Figure 1: Overlay of 1,000 samples from the rejection method (orange) and 1,000 samples from the direct method (blue) drawn from the target distribution  $f(x) = \frac{3}{8}(1 + x^2)$ , shown by the red curve.

## Question 1: Random Numbers from a Non-Flat Distribution

This question focuses on the aspect of random numbers from a non-flat distribution, namely those obeying the commonly used non-isotropic scattering function,

$$f(x) = P(\mu)d\mu = \frac{3}{8}(1 + \mu^2)d\mu \quad (1)$$

where  $\mu = \cos \theta$ .

A linear congruential generator (LCG), based on sequences of the following form was implemented as a starting point to generate pseudo-random numbers,

$$X_n = (aX_{n-1} + c) \bmod m$$

where  $a, c$  and  $m$  are integer numbers.

The rejection method generates candidate  $x$  values uniformly scaled to the interval  $[-1, 1]$  and candidate  $y$  values scaled to the range  $[0, 0.75]$  using the LCG for each iteration. For each candidate  $x$ , the value of the target function,  $f(x)$ , is computed using Eq.1, and the candidate is accepted if the corresponding  $y$  value is less than or equal to  $f(x)$ .

In contrast, the direct method employs the inverse transform technique. For a random variable with cumulative distribution function (CDF)  $F(x)$ , if  $U$  is uniformly distributed in  $[0, 1]$ , then  $x = F^{-1}(U)$  is distributed according to the PDF. The CDF is computed by integrating Eq.1, from  $-1$  to  $x$ . Although the integral leads to a cubic equation in  $x$ , it can be inverted in closed form using methods from algebra. A uniformly distributed random number  $U$  is used to directly compute a sample  $x$  by inverting the cumulative distribution function, which involves solving a cubic equation through the use of cube-root and square-root functions. This approach guarantees that each iteration produces a valid sample, eliminating the need for any rejection. To further enhance computational efficiency, a lookup table is precomputed that maps a finely spaced grid of  $U$  values to their corresponding  $x$  values. At runtime, instead of repeatedly solving the cubic equation, the algorithm uses fast linear interpolation on this lookup table to approximate  $x$  for any given  $U$ . This pre-computation significantly reduces the computational overhead while maintaining accuracy, thereby speeding up the direct method considerably.

Both methods are timed using the system clock, and their outputs are saved to separate files, providing a practical demonstration of both the theoretical and computational aspects of non-isotropic random number generation in Monte Carlo simulations. The use of separate seed variables ensures the sequences of random numbers used in each part are independent. This avoids accidental correlations between  $x$  and  $y$  coordinates and between the rejection and direct methods.




Figure 2: Independent scatter plots showing 10,000 samples each for the rejection method (top, orange points) and the direct method (bottom, blue points), drawn from the distribution. The red curve depicts the probability density function as previously shown.

Fig.1 confirms how each sampling method populates the region under the curve, ultimately yielding the desired outcome and approaching the distribution. The figure was produced with a target of 1,000 samples. A simulation producing 10,000 samples was generated and is shown in Fig.2, showing that even with a higher density of points, both methods produce results below the distribution function. Overall, as the sample size increases from 1,000 to 10,000, the shape of  $f(x)$  becomes more clearly outlined by the density of the points, demonstrating that both methods converge to the correct distribution.

To supplement the comparison between the rejection and direct sampling methods, Fig.3 was produced to show the computation times per  $N$  value of samples averaged over 10 runs. In addition, the ratio in computation time denoted as *speed-up* is shown, demonstrating the direct method's superior advantages when it comes to processing time, with an average increase in the processing time of 2.49 times for the samples plotted.

## Question 2: Monte Carlo Scattering, Isotropic

In this section, the random walk approach is used to simulate the Monte Carlo radiative transfer of photons through an atmosphere. The atmospheric region is defined by the vertical boundaries

$$z_{\min} = 0 \quad \text{and} \quad z_{\max} = 200,$$

with an overall optical depth,  $\tau$ , of

$$\tau_{\text{total}} = 10.$$

The conversion from optical depth to physical distance is handled via the parameter alpha,

$$\alpha = \frac{\tau_{\text{total}}}{z_{\max} - z_{\min}} = \frac{10}{200}.$$

This parameter is used to scale the random free path, which is sampled from an exponential distribution. The simulation proceeds as follows:

1. Each photon is launched from the bottom of the atmosphere (at  $z = 0$ ) with an initial position of  $(x, y, z) = (0, 0, 0)$ . Its initial direction is chosen isotropically. Two random numbers are used: one to sample the azimuthal angle  $\phi$  (via  $\phi = 2\pi r_1$ ) and another to determine the cosine of the polar angle ( $\mu = 2r_2 - 1$ ). The polar angle  $\theta$  is then derived from  $\mu$ , with the corresponding sine calculated as  $\sin \theta = \sqrt{1 - \mu^2}$ . The initial direction vector is given by

$$\text{dir}_x = \sin \theta \cos \phi, \quad \text{dir}_y = \sin \theta \sin \phi, \quad \text{dir}_z = \mu.$$




Figure 3: Comparison of the computation time (top panel) for the Rejection Method (orange) and the Direct Method (blue) over a range of sample sizes, plotted on a log scale. The lower panel shows the speed-up ratio (Rejection time / Direct time). Although the Rejection Method initially competes well, the Direct Method’s performance becomes more advantageous at larger sample counts, as reflected by the rising speed-up.

2. Within the atmospheric region, the photon propagates by taking steps of random length. The step length is determined by sampling an optical depth  $\tau_{\text{step}}$  from the exponential distribution  $P(\tau) = e^{-\tau}$  using the transformation

$$\tau_{\text{step}} = -\ln(U),$$

where  $U$  is a uniformly distributed random number. The physical distance  $s$  travelled in one step is then computed as:

$$s = \frac{\tau_{\text{step}}}{\alpha}.$$

The photon’s position is updated using

$$x += s \text{ dir}_x, \quad y += s \text{ dir}_y, \quad z += s \text{ dir}_z.$$

3. After each step, the code checks whether the photon has exited the atmospheric region:
  - If  $z > z_{\text{max}}$ , the photon has escaped from the top. The exiting angle is determined from the current direction (specifically, the  $z$ -component, which is  $\mu_{\text{exit}} = \text{dir}_z$ ), and this value is binned into one of 10 intervals between 0 and 1.
  - If  $z < z_{\text{min}}$ , the photon has escaped from the bottom, and it is discarded.
  - If the photon is still within the atmospheric region, a new random number is generated to decide whether it scatters or is absorbed. With an albedo of 1 (i.e. no absorption), scattering always occurs. In this case, a new direction is sampled in the same manner as at launch, giving a new set of angles ( $\phi$  and  $\theta$ ) that reset the photon’s propagation direction. This is depicted in Fig.4
4. The simulation iterates over incoming photons until one million photons have been observed to escape from the top of the atmospheric region. During this process, the code tracks the total number of photons launched, the number escaping from the top, and those escaping from the bottom.




Figure 4: A schematic of the Monte Carlo random-walk approach used to model photon transport through an atmospheric region of height  $z = 200$ . Photons are launched at  $z = 0$ , scatter randomly (black dots), and may exit the top at various angles  $\theta$ . Each dot represents a scattering event where the photon's direction is re-randomised.

5. Finally, the fraction of photons escaping in each of the 10 angular bins (based on  $\mu$ ) is calculated. The binning process involves determining the midpoint for each bin, and the resulting data is written to an output file. The final output includes the bins and the normalised intensity for each.

In addition, OpenMP is used to run many photon trajectories in parallel, thereby speeding up the Monte Carlo simulation. Each thread processes photons in batches using constant `CHUNK_SIZE`, rather than updating global counters on every photon. Locally, each thread accumulates values in counters such as `local_top`, `local_bottom`, and `local_launched`, which track how many photons escaped the top, escaped the bottom, or were launched by that thread. Periodically, these local totals are merged into the shared counters `top_escaped_global`, `bottom_escaped_global`, and `photons_launched_global` through an atomic capture block. Each thread also maintains its own random-number seeds (`local_seedA`, `local_seedB`) so random walks run independently in parallel. The code checks if `top_escaped_global` has reached one million photons, and if so, the threads stop. By structuring the code this way, using thread-local counters plus periodic atomic updates, the simulation efficiently distributes the workload across CPU cores and avoids excessive synchronisation, greatly reducing the runtime needed to achieve one million top-escaped photons.

Fig.5 shows how the normalised intensity of photons escaping through the top boundary depends on  $\cos \theta$ , where  $\theta$  is the polar angle measured from the atmospheric normal, depicted in Fig.4. Due to the atmosphere being optically thick,  $\tau_{\text{total}} = 10$ , photons directed more vertically (higher  $\cos \theta$ ) have a shorter path to exit and thus a greater likelihood of escaping. As a result, the intensity is higher at large  $\cos \theta$  values. In contrast, photons travelling at steeper angles (smaller  $\cos \theta$ ) must traverse a longer path within the atmospheric region and are more likely to scatter multiple times or exit through the bottom, leading to lower escape fractions at smaller  $\cos \theta$ . This trend is clearly reflected in the increasing curve from left to right in the plot.

### Question 3: Rayleigh scattering

The final section is concerned with scattering in the non-isotropic case, Rayleigh scattering. Physically, Rayleigh scattering describes the probability that a photon scatters through an angle  $\hat{\theta}$  between its *incoming* and *outgoing* directions. The corresponding differential cross-section can be written as

$$P(\hat{\theta}) d\Omega = \frac{3}{4} (1 + \cos^2 \hat{\theta}) \frac{d\Omega}{4\pi},$$




Figure 5: The fraction of photons escaping through the top of the atmospheric region, binned by  $\cos \theta$ . Higher values of  $\cos \theta$  (more vertical angles) yield greater escape fractions, reflecting the shorter path through the atmospheric region.

where  $\hat{\theta}$  is the angle between the incident and scattered rays, and  $d\Omega$  is an element of solid angle.

This question builds on the code in question 2, utilising the `simulate_photon_transport()` function which models the random walk of photons through the atmosphere. Unlike the isotropic case, proper Rayleigh scattering requires that the scattering angle be measured relative to the photon's incoming direction. To achieve this, the scattering angle is sampled, and then the old direction is rotated into a new direction using a local orthonormal basis. To incorporate this, several adaptations are implemented. Firstly, the absorption check remains the same as before. The code below explains the changes implemented.

#### 1. Sampling Scattering Angle

- The cosine of the scattering angle ( $\mu_s$ ) is sampled using a uniform random number passed to the `direct_sample()` function.
- This function inverts the CDF for the Rayleigh-like phase function  $f(\mu) = \frac{3}{8}(1 + \mu^2)$  to obtain  $\mu_s$ .
- Additionally, an independent random number is used to sample the azimuthal angle ( $\phi_s$ ) uniformly over  $[0, 2\pi]$ .

#### 2. Local Rotation

- A local rotation is performed of the photon's direction relative to its incoming direction.
- The current photon direction,  $\mathbf{d}_{\text{in}} = (\text{dir}_x, \text{dir}_y, \text{dir}_z)$ , is normalised to create a unit vector.
- An orthonormal basis is constructed. A vector  $\mathbf{h}$  is computed by taking the cross product of  $\mathbf{n}$  (a unit vector of the photons incoming direction) with a reference vector (usually the z-axis unless  $\mathbf{n}$  is nearly parallel to it, in which case the x-axis is used). A second perpendicular vector  $\mathbf{u}$  is then obtained by computing the cross product of  $\mathbf{n}$  and  $\mathbf{h}$ .
- New photon direction is computed using  $\mathbf{d}_{\text{out}} = \mu_s \mathbf{n} + \sqrt{1 - \mu_s^2} \cos \phi_s \mathbf{h} + \sqrt{1 - \mu_s^2} \sin \phi_s \mathbf{u}$ . This is now the direction of the photon relative to the incoming direction.

#### 3. Update photon's direction

- The new direction vector from the local rotation replaces the old direction and is used for subsequent steps in the photon's random walk.

Fig.6 demonstrates the case where Rayleigh scattering has been implemented for  $\tau = 10$  and  $\tau = 0.1$ . When photons of shorter wavelengths (blue light) propagate through the atmosphere, they experience a higher optical depth and thus scatter more frequently than longer-wavelength (redder)




Figure 6: Comparison of the normalised intensity of photons escaping from the top of the atmosphere, binned by  $\cos \theta$ , for two different optical depths of  $\tau = 10$  and  $\tau = 0.1$ . Higher  $\tau$  causes more scattering events, concentrating escaping photons near  $\cos \theta \approx 1$ , while lower  $\tau$  yields a flatter distribution over angle.




Figure 7: Comparison of the normalised intensity of photons escaping from the top of the atmosphere, binned by  $\cos \theta$ , for multiple optical depths ( $\tau = 10, 2, 1, 0.5, 0.1, 0.01$ ). Higher  $\tau$  causes more scattering events, concentrating escaping photons near  $\cos \theta = 1$ , while lower  $\tau$  yields a flatter distribution.



photons. This is modelled by assigning a larger  $\tau = 10$  for blue photons and a smaller  $\tau = 0.1$  for other colours. The higher  $\tau$  (blue) photons are scattered strongly and emerge at a wide range of angles away from the direct beam, leading to a pronounced intensity of blue light when looking overhead. Meanwhile, the direct sunlight beam loses much of its short-wavelength component, making the transmitted Sun itself appear redder, especially at lower angles. Consequently, this difference in scattering strength quantitatively explains why the sky appears blue (dominated by scattered short-wavelength photons) when one is not staring directly at the Sun.

Around  $\cos(\theta) \approx 0.9$ , the photons are travelling fairly close to vertically but not perfectly so, and in that regime, the fraction of photons that escape happens to converge for a range of optical depths. In an optically thick medium, angles near the vertical are most favoured for escape, whereas in an optically thin medium, nearly all angles have a decent escape probability. There can be a specific “pivot” angle (around  $\cos(\theta) \approx 0.9$  in Fig.7) where these effects balance out, causing curves for different  $\tau$  values to intersect. Numerics and finite sampling can further accentuate the intersection, but physically, it reflects a crossover point where neither extremely high nor extremely low optical depth dominates the escape probability.

# Appendix

## 4.1 Full Code

```
1 // -----
2 // To compile: (Capital letter O not zero)
3 // gcc -fopenmp -O3 -ffast-math -march=native 24367.c -o 24367 -lm
4 // To run:
5 // ./24367
6 // -----
7 #include <math.h>
8 #include <stdlib.h>
9 #include <stdio.h>
10 #include <time.h>
11 #include <omp.h>
12
13 // Global constants
14 const long long MODULUS = 2147483647LL; //  $2^{31} - 1$ 
15 long long seed_scatter = 99991; // For Q2 & Q3 scattering
16 long long seed1 = 1, seed2 = 107, seed3 = 123;
17
18 // Constants for Q1
19 const int NUMSAMPLES = 10000;
20 int num_accepted = 0, count = 0;
21
22 // Constants for Q2 and Q3
23 const double Z_MIN = 0.0, Z_MAX = 200.0, ALBEDO = 1.0;
24 const long long NPHOTONS_TARGET = 1000000LL; // Exactly 1,000,000 out top is needed
25 const int NUM_BINS = 10;
26 double tau_total = 10.0; // Change if needed to 0.1 for 'other colours' case
27
28 // Constants for lookup table for direct method
29 #define LUT_SIZE 10000 // Number of table entries
30 static double lookup_table[LUT_SIZE];
31 static int lookup_init = 0;
32
33 // Linear Congruential Generator (LCG) for [0,1) random
34 long long LCG(long long current, long long a, long long c, long long m) {
35     return (a * current + c) % m;
36 }
37
38 // This inline function advances the LCG-based seed and returns a random number in
39 // [0, 1).
40 static inline double LCG_rand01(long long *seedptr, long long a) {
41     *seedptr = LCG(*seedptr, a, 0, MODULUS);
42     return (double)(*seedptr) / (double)MODULUS;
43 }
44
45 // The PDF for Q1  $\frac{3}{8}(1+x^2)$  for x in [-1,1]
46 double distribution_func(double x) {
47     return (3.0 / 8.0) * (1.0 + x*x);
48 }
49
50 // Solve cubic for x in [-1,1] that satisfies F(x) = U. Inverts the CDF
51 double direct_sample(double U) {
52     double half_q = 2.0 - 4.0 * U;
53     double term = sqrt(half_q*half_q + 1.0);
54     double c1 = cbrt(-half_q + term);
55     double c2 = cbrt(-half_q - term);
56     return (c1 + c2);
57 }
58
59 // Precompute the lookup table.
60 void init_lookup_table(void) {
61     if (lookup_init)
62         return;
63     for (int i = 0; i < LUT_SIZE; i++) {
64         double U = (double)i / (LUT_SIZE - 1); // U spans 0 to 1
65         lookup_table[i] = direct_sample(U);
66     }
67     lookup_init = 1;
68 }
69
70 // Use linear interpolation on the lookup table to approximate direct_sample(U).
```

```

70 double direct_sample_lookup(double U) {
71     if (!lookup_init)
72         init_lookup_table();
73     double pos = U * (LUT_SIZE - 1);
74     int index = (int) pos;
75     double frac = pos - index;
76     if (index >= LUT_SIZE - 1)
77         return lookup_table[LUT_SIZE - 1];
78     return lookup_table[index]*(1 - frac) + lookup_table[index + 1]*frac;
79 }
80
81 // Q1: Rejection & Direct Methods (left single-threaded)
82 int question_1(void) {
83     // Allocate arrays for storing sampled points (x,y) in rejection method
84     double *rej_x = (double *)malloc(NUMSAMPLES * sizeof(double));
85     double *rej_y = (double *)malloc(NUMSAMPLES * sizeof(double));
86     if (!rej_x || !rej_y) {
87         perror("Error: Memory allocation failed");
88         return 1;
89     }
90     // — Rejection Method —
91     clock_t start_rejection = clock();
92     while (num_accepted < NUMSAMPLES) {
93         count++; // Total iterations accepted
94         // Sample x in [-1,1], y in [0,0.75]
95         double x_reject = 2.0 * LCG_rand01(&seed1, 16807) - 1.0;
96         double y_reject = 0.75 * LCG_rand01(&seed2, 48271);
97         // Evaluate PDF at x
98         double fx = distribution_func(x_reject);
99         // If (x_reject, y_reject) lies under curve then accept
100         if (y_reject <= fx) {
101             rej_x[num_accepted] = x_reject;
102             rej_y[num_accepted] = y_reject;
103             num_accepted++;
104         }
105     }
106     clock_t end_rejection = clock();
107     double time_rejection = (double)(end_rejection - start_rejection) /
CLOCKS.PER_SEC;
108     // Write rejection samples
109     FILE *fptr_rejection = fopen("rejection_method.txt", "w");
110     if (fptr_rejection) {
111         for (int i = 0; i < NUMSAMPLES; i++) {
112             fprintf(fptr_rejection, "%.3f, %.3f\n", rej_x[i], rej_y[i]);
113         }
114         fclose(fptr_rejection);
115     }
116     // — Direct Method —
117     // Allocate arrays for direct method samples
118     double *dir_x = (double *)malloc(NUMSAMPLES * sizeof(double));
119     double *dir_y = (double *)malloc(NUMSAMPLES * sizeof(double));
120     if (!dir_x || !dir_y) {
121         perror("Error: Memory allocation failed");
122         free(rej_x);
123         free(rej_y);
124         return 1;
125     }
126     // Measure the lookup table build time
127     clock_t start_lut = clock();
128     init_lookup_table();
129     clock_t end_lut = clock();
130     double time_lut = (double)(end_lut - start_lut) / CLOCKS_PER_SEC;
131     printf("Lookup table build time: %.8f s\n", time_lut);
132
133     clock_t start_direct = clock();
134     for (int i = 0; i < NUMSAMPLES; i++) {
135         // Generate U in [0,1], invert CDF to get x
136         double U = LCG_rand01(&seed3, 16807);
137         // Use lookup table with linear interpolation to compute x
138         double x_val = direct_sample_lookup(U);
139         // For plotting only: pick y as a random fraction up to f(x_val)
140         double y_val = LCG_rand01(&seed3, 48271) * distribution_func(x_val);
141         dir_x[i] = x_val;
142         dir_y[i] = y_val;
143     }
144     clock_t end_direct = clock();

```

```

145     double time_direct = (double)(end_direct - start_direct) / CLOCKS_PER_SEC;
146     // Write direct method output to file
147     FILE *fptr_direct = fopen("direct.method.txt", "w");
148     if (fptr_direct) {
149         for (int i = 0; i < NUMSAMPLES; i++) {
150             fprintf(fptr_direct, "%.3f, %.3f\n", dir_x[i], dir_y[i]);
151         }
152         fclose(fptr_direct);
153     }
154     // Print summary of results
155     printf("\nPart 1: Rejection Method:\n");
156     printf("    Accepted samples = %d\n", NUMSAMPLES);
157     printf("    Total iterations = %d\n", count);
158     printf("    Acceptance ratio = %.3f\n", (double)NUMSAMPLES / (double)count);
159     printf("    Total CPU time (s)= %.8f\n", time_rejection);
160     printf("    Time per sample = %.8e s\n\n", time_rejection / NUMSAMPLES);
161
162     printf("Part 1: Direct Method (Lookup Table):\n");
163     printf("    Generated samples = %d\n", NUMSAMPLES);
164     printf("    Total CPU time (s)= %.8f\n", time_direct);
165     printf("    Time per sample = %.8e s\n\n", time_direct / NUMSAMPLES);
166     // Avoid dividing by zero if times are extremely small
167     if (time_rejection <= 0.0) time_rejection = 1e-6;
168     if (time_direct <= 0.0) time_direct = 1e-6;
169     double speedup = time_rejection / time_direct;
170     printf("Speedup (Rejection / Direct) = %.4f\n\n", speedup);
171     // Free all allocated memory
172     free(rej_x); free(rej_y);
173     free(dir_x); free(dir_y);
174     return 0;
175 }
176
177 // — Isotropic Scattering for Q2 —
178 double scatter_isotropic(double U) {
179     // Map uniform random U in [0,1) to mu in [-1,1]
180     // for an isotropic phase function: mu = 2U - 1
181     return 2.0*U - 1.0;
182 }
183
184 // — Rayleigh scattering for Q3 —
185 double scatter_rayleigh(double U) {
186     // Rayleigh scattering ~ (1 + cos^2 theta),
187     // reuse direct_sample(U) which inverts the CDF.
188     return direct_sample(U);
189 }
190
191 // write_bins: outputs bin index, fraction, mu_mid to 'fout'
192 static void write_bins(int *bin_counts, int num_bins, double dmu, long long
193     top_escaped, FILE *fout) {
194     for (int j = 0; j < num_bins; j++) {
195         // mu_mid is the midpoint for the j-th bin in [0,1]
196         double mu_mid = (j + 0.5) * dmu;
197         // fraction = fraction of total top-escaped photons in this bin
198         double fraction = 0.0;
199         if (top_escaped > 0) {
200             fraction = (double)bin_counts[j] / (double)top_escaped;
201         }
202         fprintf(fout, "%d, %e, %f\n", j, fraction, mu_mid);
203     }
204 }
205
206 // Orthonormal rotation for Rayleigh scattering in local frame.
207 // rotate the direction vector 'd_in' into 'd_out' with a scattering angle mu and
208 // azimuth phi.
209 static void rotate_rayleigh_local(
210     const double d_in[3], // Old direction
211     double mu, // cos (scattering angle)
212     double phi, // random azimuth in [0,2 )
213     double d_out[3]) // New direction
214 {
215     // Normalise the old direction
216     double len = sqrt(d_in[0]*d_in[0] + d_in[1]*d_in[1] + d_in[2]*d_in[2]);
217     double nx = d_in[0]/len, ny = d_in[1]/len, nz = d_in[2]/len;
218     double sin_theta = sqrt(1.0 - mu*mu);
219
220     // Build local orthonormal basis

```

```

219 // Find a vector h orthonormal to n then cross to get the 3rd vector u
220 double hx, hy, hz;
221 if (fabs(nz) < 0.9999) {
222     // cross n with z-hat -> h = (ny, -nx, 0)
223     hx = ny;
224     hy = -nx;
225     hz = 0.0;
226 } else {
227     // cross n with x-hat -> h = (0, nz, -ny)
228     hx = 0.0;
229     hy = nz;
230     hz = -ny;
231 }
232 double h_len = sqrt(hx*hx + hy*hy + hz*hz);
233 hx /= h_len;
234 hy /= h_len;
235 hz /= h_len;
236
237 // Now u = n x h
238 double ux = ny*hz - nz*hy;
239 double uy = nz*hx - nx*hz;
240 double uz = nx*hy - ny*hx;
241
242 // The new direction in local coordinates
243 d_out[0] = mu*nx + sin_theta*cos(phi)*hx + sin_theta*sin(phi)*ux;
244 d_out[1] = mu*ny + sin_theta*cos(phi)*hy + sin_theta*sin(phi)*uy;
245 d_out[2] = mu*nz + sin_theta*cos(phi)*hz + sin_theta*sin(phi)*uz;
246 }
247
248 // Q2 and Q3 driver function to determine scattering type
249 typedef enum {
250     INIT_ISOTROPIC = 0,
251     INIT_VERTICAL = 1
252 } InitDirectionType;
253
254 // Function pointer for scattering routines
255 typedef double (*ScatteringFunc)(double);
256
257 // Launch photons until n_photons_required escape the top boundary.
258 // Each photon is launched (either isotropic or vertical initial dir),
259 // and scatters (Rayleigh or Isotropic). Final directions are binned by mu in 10 bins
260
261 // If top_escaped >= n_photons_required, the simulation stops.
262 void simulate_photon_transport(
263     double tau_total, double zmin, double zmax,
264     double albedo, long long n_photons_required,
265     InitDirectionType init_type, ScatteringFunc scatter_func,
266     const char *outfile_label) {
267     // dmu = bin width in mu from 0..1 for 10 bins
268     double dmu = 1.0 / (double)NUMBINS;
269     int *bin_counts_global = (int *)calloc(NUMBINS, sizeof(int));
270     if (!bin_counts_global) {
271         fprintf(stderr, "Error: Could not allocate bin_counts.\n");
272         return;
273     }
274     FILE *fout = fopen(outfile_label, "w");
275     if (!fout) {
276         perror("Error opening output file");
277         free(bin_counts_global);
278         return;
279     }
280     // alpha = tau_total / (zmax - zmin)
281     double alpha = tau_total / (zmax - zmin);
282
283     static long long top_escaped_global = 0;
284     static long long bottom_escaped_global = 0;
285     static long long photons_launched_global = 0;
286     // Make sure each run starts from zero
287     top_escaped_global = 0;
288     bottom_escaped_global = 0;
289     photons_launched_global = 0;
290
291     // Process in chunks
292     const int CHUNK_SIZE = 1000;
293     // Start timing the simulation
294     clock_t start_time = clock();

```

```

294
295 // Start parallel region
296 #pragma omp parallel
297 {
298     // Each thread has local seeds, counters
299     int tid = omp_get_thread_num();
300     long long local_seedA = seed_scatter + 10000LL * tid;
301     long long local_seedB = (seed_scatter + 12345LL) + 10000LL * tid;
302
303     long long local_top = 0; // how many escaped top from this thread
304     long long local_bottom = 0; // how many escaped bottom
305     long long local_launched = 0;
306
307     int bin_counts_local[NUM_BINS];
308     for (int i = 0; i < NUM_BINS; i++) {
309         bin_counts_local[i] = 0;
310     }
311
312     // Keep going until top_escaped_global >= n_photons_required
313     while (1) {
314         long long curr_top;
315         #pragma omp atomic read
316         curr_top = top_escaped_global;
317         if (curr_top >= n_photons_required) {
318             break;
319         }
320
321         long long needed = n_photons_required - curr_top;
322         long long chunk = (needed < CHUNK_SIZE) ? needed : CHUNK_SIZE;
323
324         // Launch chunk photons
325         for (int c = 0; c < chunk; c++) {
326             local_launched++;
327
328             // Initialise photon position and direction
329             double x=0.0, y=0.0, z=0.0;
330             double phi, mu;
331             // If init_type == INIT_ISOTROPIC -> random direction
332             if (init_type == INIT_ISOTROPIC) {
333                 double r1 = LCG_rand01(&local_seedA, 16807);
334                 double r2 = LCG_rand01(&local_seedB, 48271);
335                 phi = 2.0 * M_PI * r1;
336                 mu = 2.0 * r2 - 1.0;
337             // else (INIT_VERTICAL) -> mu=1, phi=0 -> straight up
338             } else {
339                 mu = 1.0;
340                 phi = 0.0;
341             }
342             double sin_theta = sqrt(1.0 - mu*mu);
343             double dir_x = sin_theta*cos(phi);
344             double dir_y = sin_theta*sin(phi);
345             double dir_z = mu;
346
347             // Random walk
348             while (1) {
349                 double tau_step = -log(LCG_rand01(&local_seedA, 16807));
350                 double s = tau_step / alpha;
351                 x += s*dir_x;
352                 y += s*dir_y;
353                 z += s*dir_z;
354                 // If photon escapes top, bin by mu
355                 if (z > zmax) {
356                     double mu_exit = dir_z;
357                     if (mu_exit < 0.0) mu_exit = 0.0;
358                     int index = (int)(mu_exit / dmu);
359                     if (index >= NUM_BINS) index = NUM_BINS - 1;
360                     bin_counts_local[index]++;
361                     local_top++;
362                     break;
363                 } else if (z < zmin) {
364                     // If it escapes bottom
365                     local_bottom++;
366                     break;
367                 } else {
368                     // Possibly scatter or absorb
369                     double rscat = LCG_rand01(&local_seedB, 48271);

```

```

370         if (rscat < albedo) {
371             // Decide new direction
372             if (scatter_func == scatter_rayleigh) {
373                 double rU = LCG_rand01(&local_seedA, 16807);
374                 double mu_s = direct_sample(rU);
375                 double r_phi = LCG_rand01(&local_seedB, 48271);
376                 double phi_s = 2.0 * M_PI * r_phi;
377
378                 double d_in[3] = {dir_x, dir_y, dir_z};
379                 double d_out[3];
380                 // rotate around old direction using mu_s, phi_s
381                 rotate_rayleigh_local(d_in, mu_s, phi_s, d_out);
382                 dir_x = d_out[0];
383                 dir_y = d_out[1];
384                 dir_z = d_out[2];
385             } else {
386                 // Isotropic
387                 double r_phi = LCG_rand01(&local_seedA, 16807);
388                 double r_mu = LCG_rand01(&local_seedB, 48271);
389                 double phi_s = 2.0*M_PI * r_phi;
390                 double mu_s = scatter_func(r_mu);
391
392                 double stheta = sqrt(1.0 - mu_s*mu_s);
393                 dir_x = stheta * cos(phi_s);
394                 dir_y = stheta * sin(phi_s);
395                 dir_z = mu_s;
396             }
397         } else {
398             // absorbed -> done
399             break;
400         }
401     }
402 } // end random walk
403 } // end chunk loop from parralel region
404
405 // Merge local counters into global with atomic capture
406 long long old_top, new_top;
407 #pragma omp atomic capture
408 {
409     old_top = top_escaped_global;
410     top_escaped_global = top_escaped_global + local_top;
411 }
412 new_top = old_top + local_top;
413 if (new_top > n_photons_required) {
414     #pragma omp critical
415     {
416         if (top_escaped_global > n_photons_required) {
417             top_escaped_global = n_photons_required;
418         }
419     }
420 }
421 // reset local_top so it is not added it again
422 local_top = 0;
423 // update bottom
424 #pragma omp atomic
425 bottom_escaped_global += local_bottom;
426 local_bottom = 0;
427 // update launched
428 #pragma omp atomic
429 photons_launched_global += local_launched;
430 local_launched = 0;
431 // merge bin arrays
432 #pragma omp critical
433 {
434     for (int i = 0; i < NUM_BINS; i++) {
435         bin_counts_global[i] += bin_counts_local[i];
436         bin_counts_local[i] = 0;
437     }
438 }
439 // If at or above the target, break
440 long long after_top;
441 #pragma omp atomic read
442 after_top = top_escaped_global;
443 if (after_top >= n_photons_required) {
444     break;
445 }

```

```

446     } // end while loop
447 } // end parallel region
448
449 clock_t end_time = clock();
450 double elapsed = (double)(end_time - start_time) / CLOCKS_PER_SEC;
451
452 printf("Simulation: %s\n", outfile_label);
453 printf("    Photons launched : %lld\n", photons_launched_global);
454 printf("    Escaped top      : %lld\n", top_escaped_global);
455 printf("    Escaped bottom    : %lld\n", bottom_escaped_global);
456 printf("    CPU time (s)      : %.8f\n\n", elapsed);
457
458 // Write final bin data to file
459 write_bins(bin_counts_global, NUM_BINS, dm_u, top_escaped_global, fout);
460 fclose(fout);
461 free(bin_counts_global);
462 }
463
464 // Q2: Isotropic scattering
465 int question_2(void) {
466     simulate_photon_transport(
467         tau_total, Z_MIN, Z_MAX, ALBEDO,
468         NPHOTONS_TARGET, INIT_ISOTROPIC, scatter_isotropic,
469         "Question_2_Isotropic.txt");
470     return 0;
471 }
472
473 // Q3: Rayleigh scattering
474 int question_3(void) {
475     simulate_photon_transport(
476         tau_total, Z_MIN, Z_MAX, ALBEDO,
477         NPHOTONS_TARGET, INIT_VERTICAL, scatter_rayleigh,
478         "Question_3_Rayleigh.txt");
479     return 0;
480 }
481
482 int main(void) {
483     question_1();
484     question_2();
485     question_3();
486     return 0;
487 }

```

**End of Report**